

## Task 2

### Credit / Home Loans - AutoML vs Bespoke ML

Standard Bank is embracing the digital transformation wave and intends to use new and exciting technologies to give their customers a complete set of services from the convenience of their mobile devices. As Africa's biggest lender by assets, the bank aims to improve the current process in which potential borrowers apply for a home loan. The current process involves loan officers having to manually process home loan applications. This process takes 2 to 3 days to process upon which the applicant will receive communication on whether or not they have been granted the loan for the requested amount. To improve the process Standard Bank wants to make use of machine learning to assess the credit worthiness of an applicant by implementing a model that will predict if the potential borrower will default on his/her loan or not, and do this such that the applicant receives a response immediately after completing their application.

You will be required to follow the data science lifecycle to fulfill the objective. The data science lifecycle (<https://www.datascience-pm.com/crisp-dm-2/>) includes:

- Business Understanding
- Data Understanding
- Data Preparation
- Modelling
- Evaluation
- Deployment.

You now know the Cross Industry Standard Process for Data Mining (CRISP-DM), have an idea of the business needs and objectivess, and understand the data. Next is the tedious task of preparing the data for modeling, modeling and evaluating the model. Luckily, just like EDA the first of the two phases can be automated. But also, just like EDA this is not always best.

In this task you will be get a taste of AutoML and Bespoke ML. In the notebook we make use of the library `auto-sklearn` (<https://www.automl.org/automl/auto-sklearn/>) for AutoML and `sklearn` for ML. We will use train one machine for the traditional approach and you will be required to change this model to any of the models that exist in `sklearn`. The model we will train will be a Logistic Regression. Parts of the data preparation will be omitted for you to do, but we will provide hints to lead you in the right direction.

The data provided can be found in the Resources folder as well as (<https://www.kaggle.com/datasets/altruistdelhite04/loan-prediction-problem-dataset>).

- train will serve as the historical dataset that the model will be trained on and,
- test will serve as unseen data we will predict on, i.e. new ('future') applicants.

### Part One

There are many AutoEDA Python libraries out there which include:

- dtale (<https://dtale.readthedocs.io/en/latest/>)
- pandas profiling (<https://pandas-profiling.ydata.ai/docs/master/index.html>)
- autoviz (<https://readthedocs.org/projects/autoviz/>)
- sweetviz (<https://pypi.org/project/sweetviz/>)

and many more. In this task we will use Sweetviz.. You may be required to use bespoke EDA methods.

The Home Loans Department manager wants to know the following:

1. An overview of the data. (HINT: Provide the number of records, fields and their data types. Do for both).
2. What data quality issues exist in both train and test? (HINT: Comment any missing values and duplicates)
3. How do the the loan statuses compare? i.e. what is the distrubition of each?
4. How many of the loan applicants have dependents based on the historical dataset?
5. How do the incomes of those who are employed compare to those who are self employed based on the historical dataset?

historical dataset:

6. Are applicants with a credit history more likely to default than those who do not have one?
7. Is there a correlation between the applicant's income and the loan amount they applied for?

## Part Two

Run the AutoML section and then fill in code for the traditional ML section for the the omitted cells.

Please note that the notebook you submit must include the analysis you did in Task 2.

## Import Libraries

In [2]:

```
!pip install tpot
```

```
Requirement already satisfied: tpot in c:\users\rhydh\anaconda3\lib\site-packages (0.12.2)
Requirement already satisfied: numpy>=1.16.3 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (1.26.4)
Requirement already satisfied: scipy>=1.3.1 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (1.13.1)
Requirement already satisfied: scikit-learn>=1.4.1 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (1.5.0)
Requirement already satisfied: deap>=1.2 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (1.4.1)
Requirement already satisfied: update-checker>=0.16 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (0.18.0)
Requirement already satisfied: tqdm>=4.36.1 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (4.65.0)
Requirement already satisfied: stopit>=1.1.1 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (1.1.2)
Requirement already satisfied: pandas>=0.24.2 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (2.1.4)
Requirement already satisfied: joblib>=0.13.2 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (1.2.0)
Requirement already satisfied: xgboost>=1.1.0 in c:\users\rhydh\anaconda3\lib\site-packages (from tpot) (2.0.3)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\rhydh\anaconda3\lib\site-packages (from pandas>=0.24.2->tpot) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in c:\users\rhydh\anaconda3\lib\site-packages (from pandas>=0.24.2->tpot) (2023.3.post1)
Requirement already satisfied: tzdata>=2022.1 in c:\users\rhydh\anaconda3\lib\site-packages (from pandas>=0.24.2->tpot) (2023.3)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\rhydh\anaconda3\lib\site-packages (from scikit-learn>=1.4.1->tpot) (3.5.0)
Requirement already satisfied: colorama in c:\users\rhydh\anaconda3\lib\site-packages (from tqdm>=4.36.1->tpot) (0.4.6)
Requirement already satisfied: requests>=2.3.0 in c:\users\rhydh\anaconda3\lib\site-packages (from update-checker>=0.16->tpot) (2.31.0)
Requirement already satisfied: six>=1.5 in c:\users\rhydh\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas>=0.24.2->tpot) (1.16.0)
Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\rhydh\anaconda3\lib\site-packages (from requests>=2.3.0->update-checker>=0.16->tpot) (2.0.4)
Requirement already satisfied: idna<4,>=2.5 in c:\users\rhydh\anaconda3\lib\site-packages (from requests>=2.3.0->update-checker>=0.16->tpot) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\rhydh\anaconda3\lib\site-packages (from requests>=2.3.0->update-checker>=0.16->tpot) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\rhydh\anaconda3\lib\site-packages (from requests>=2.3.0->update-checker>=0.16->tpot) (2024.2.2)
```

In [3]:

```
!pip install --upgrade scipy
```

```
Requirement already satisfied: scipy in c:\users\rhydh\anaconda3\lib\site-packages (1.13.1)
Requirement already satisfied: numpy<2.3,>=1.22.4 in c:\users\rhydh\anaconda3\lib\site-packages (from scipy) (1.26.4)
```

In [4]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sweetviz
from tpot import TPOTClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.impute import SimpleImputer
```

```
C:\Users\rhydh\anaconda3\Lib\site-packages\tpot\builtins\__init__.py:36: UserWarning: Warning: optional dependency `torch` is not available. - skipping import of NN models.
  warnings.warn("Warning: optional dependency `torch` is not available. - skipping import of NN models.")
```

## Import Datasets

In [5]:

```
train = pd.read_csv('StandardBank train.csv')
test = pd.read_csv('StandardBank test.csv')
```

## Part One

### EDA

In [6]:

```
train.head()
```

Out[6]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Status
0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	

In [7]:

```
test.head()
```

Out[7]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Status
0	LP001015	Male	Yes	0	Graduate	No	5720	0	110.0	
1	LP001022	Male	Yes	1	Graduate	No	3076	1500	126.0	
2	LP001031	Male	Yes	2	Graduate	No	5000	1800	208.0	
3	LP001035	Male	Yes	2	Graduate	No	2340	2546	100.0	
4	LP001051	Male	No	0	Not Graduate	No	3276	0	78.0	

In [8]:

```
# we concat for easy analysis
n = train.shape[0] # we set this to be able to separate the
df = pd.concat([train, test], axis=0)
df.head()
```

Out[8]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Status
0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN	0
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	1
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	1
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	1
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	1

Sweetviz

In [9]:

```
autoEDA = sweetviz.analyze(train)
autoEDA.show_notebook()
```

## Your Own EDA

In [10]:

```
df.describe(include='all')
```

Out[10]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount
count	981	957	978	956	981	926	981.000000	981.000000	954.000000
unique	981	2	2	4	2	2	NaN	NaN	NaN
top	LP001002	Male	Yes	0	Graduate	No	NaN	NaN	NaN
freq	1	775	631	545	763	807	NaN	NaN	NaN
mean	NaN	NaN	NaN	NaN	NaN	NaN	5179.795107	1601.916330	142.511530
std	NaN	NaN	NaN	NaN	NaN	NaN	5695.104533	2718.772806	77.421740
min	NaN	NaN	NaN	NaN	NaN	NaN	0.000000	0.000000	9.000000
25%	NaN	NaN	NaN	NaN	NaN	NaN	2875.000000	0.000000	100.000000
50%	NaN	NaN	NaN	NaN	NaN	NaN	3800.000000	1110.000000	126.000000
75%	NaN	NaN	NaN	NaN	NaN	NaN	5516.000000	2365.000000	162.000000
max	NaN	NaN	NaN	NaN	NaN	NaN	81000.000000	41667.000000	700.000000

In [11]:

```
df.dtypes
```

Out[11]:

```
Loan_ID          object
Gender           object
Married          object
Dependents       object
Education        object
Self_Employed    object
ApplicantIncome  int64
CoapplicantIncome float64
LoanAmount       float64
Loan_Amount_Term float64
Credit_History   float64
Property_Area     object
Loan_Status      object
dtype: object
```

In [12]:

```
df.columns
```

Out[12]:

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
       'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
       'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

In [13]:

```
train.isna().sum()
```

Out[13]:

```
Loan_ID      0
Gender      13
Married      3
Dependents   15
Education    0
Self_Employed 32
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount   22
Loan_Amount_Term 14
Credit_History 50
Property_Area 0
Loan_Status  0
dtype: int64
```

In [14]:

```
test.isna().sum()
```

Out[14]:

```
Loan_ID      0
Gender      11
Married      0
Dependents   10
Education    0
Self_Employed 23
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount    5
Loan_Amount_Term 6
Credit_History 29
Property_Area  0
dtype: int64
```

In [15]:

```
# Convert float values to 1 or 0
train['Credit_History'] = train['Credit_History'].apply(lambda x: '1' if x >= 1 else '0'
)

# Convert dtype to object
train['Credit_History'] = train['Credit_History'].astype('object')

# Display the modified DataFrame
print(train)
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
..	...	...	...	...	...	...	
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	
	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\		
0	5849	0.0	NaN	360.0			
1	4583	1508.0	128.0	360.0			
2	3000	0.0	66.0	360.0			
3	2583	2358.0	120.0	360.0			
4	6000	0.0	141.0	360.0			
..	...	...	...	...			
609	2900	0.0	71.0	360.0			
610	4106	0.0	40.0	180.0			
611	8072	240.0	253.0	360.0			
612	7500	0.0	127.0	360.0			

```
612          7583          0.0          187.0          360.0
613          4583          0.0          133.0          360.0
```

```
      Credit_History Property_Area Loan_Status
0              1          Urban            Y
1              1          Rural            N
2              1          Urban            Y
3              1          Urban            Y
4              1          Urban            Y
..          ...          ...          ...
609             1          Rural            Y
610             1          Rural            Y
611             1          Urban            Y
612             1          Urban            Y
613             0      Semiurban            N
```

[614 rows x 13 columns]

In [16]:

```
# Convert float values to 1 or 0
test['Credit_History'] = test['Credit_History'].apply(lambda x: '1' if x >= 1 else '0')

# Convert dtype to object
test['Credit_History'] = test['Credit_History'].astype('object')

# Display the modified DataFrame
print(test)
```

```
      Loan_ID Gender Married Dependents      Education Self_Employed \
0    LP001015  Male    Yes           0      Graduate           No
1    LP001022  Male    Yes           1      Graduate           No
2    LP001031  Male    Yes           2      Graduate           No
3    LP001035  Male    Yes           2      Graduate           No
4    LP001051  Male    No            0  Not Graduate           No
..          ...    ...    ...          ...          ...          ...
362  LP002971  Male    Yes          3+  Not Graduate           Yes
363  LP002975  Male    Yes           0      Graduate           No
364  LP002980  Male    No            0      Graduate           No
365  LP002986  Male    Yes           0      Graduate           No
366  LP002989  Male    No            0      Graduate           Yes
```

```
      ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term \
0              5720              0          110.0          360.0
1              3076             1500          126.0          360.0
2              5000             1800          208.0          360.0
3              2340             2546          100.0          360.0
4              3276              0           78.0          360.0
..          ...          ...          ...          ...
362             4009             1777          113.0          360.0
363             4158              709          115.0          360.0
364             3250             1993          126.0          360.0
365             5000             2393          158.0          360.0
366             9200              0           98.0          180.0
```

```
      Credit_History Property_Area
0              1          Urban
1              1          Urban
2              1          Urban
3              0          Urban
4              1          Urban
..          ...          ...
362             1          Urban
363             1          Urban
364             0      Semiurban
365             1          Rural
366             1          Rural
```

[367 rows x 12 columns]

In [17]:

```

# Initialize SimpleImputer for numeric columns
numeric_imputer = SimpleImputer(strategy='mean')

# Initialize SimpleImputer for object columns
object_imputer = SimpleImputer(strategy='most_frequent')

# Loop through columns
for col in train.columns:
    if train[col].dtype == 'object':
        # Impute object columns with most frequent value
        train[col] = object_imputer.fit_transform(train[[col]])[:, 0] # Extracting the
imputed values from the 2D array
    elif train[col].dtype in ['int64', 'float64']: # Adjust as per your specific numeri
c types
        # Impute numeric columns with mean
        train[col] = numeric_imputer.fit_transform(train[[col]])[:, 0] # Extracting the
imputed values from the 2D array

# Display the imputed DataFrame
print(train)

```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
..	...	...	...	...	...	...	
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849.0	0.0	146.412162	360.0	
1	4583.0	1508.0	128.000000	360.0	
2	3000.0	0.0	66.000000	360.0	
3	2583.0	2358.0	120.000000	360.0	
4	6000.0	0.0	141.000000	360.0	
..	...	...	...	...	
609	2900.0	0.0	71.000000	360.0	
610	4106.0	0.0	40.000000	180.0	
611	8072.0	240.0	253.000000	360.0	
612	7583.0	0.0	187.000000	360.0	
613	4583.0	0.0	133.000000	360.0	

	Credit_History	Property_Area	Loan_Status
0	1	Urban	Y
1	1	Rural	N
2	1	Urban	Y
3	1	Urban	Y
4	1	Urban	Y
..	...	...	...
609	1	Rural	Y
610	1	Rural	Y
611	1	Urban	Y
612	1	Urban	Y
613	0	Semiurban	N

[614 rows x 13 columns]

In [18]:

```

# Loop through columns
for col in test.columns:
    if test[col].dtype == 'object':
        # Impute object columns with most frequent value
        test[col] = object_imputer.fit_transform(test[[col]])[:, 0] # Extracting the imp
uted values from the 2D array
    elif test[col].dtype in ['int64', 'float64']: # Adjust as per your specific numeric
types

```



```

# Impute numeric columns with mean
test[col] = numeric_imputer.fit_transform(test[[col]])[:, 0] # Extracting the imputed values from the 2D array

# Display the imputed DataFrame
print(test)

```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001015	Male	Yes	0	Graduate	No	
1	LP001022	Male	Yes	1	Graduate	No	
2	LP001031	Male	Yes	2	Graduate	No	
3	LP001035	Male	Yes	2	Graduate	No	
4	LP001051	Male	No	0	Not Graduate	No	
..	...	...	...	...	...	...	
362	LP002971	Male	Yes	3+	Not Graduate	Yes	
363	LP002975	Male	Yes	0	Graduate	No	
364	LP002980	Male	No	0	Graduate	No	
365	LP002986	Male	Yes	0	Graduate	No	
366	LP002989	Male	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5720.0	0.0	110.0	360.0	
1	3076.0	1500.0	126.0	360.0	
2	5000.0	1800.0	208.0	360.0	
3	2340.0	2546.0	100.0	360.0	
4	3276.0	0.0	78.0	360.0	
..	...	...	...	...	
362	4009.0	1777.0	113.0	360.0	
363	4158.0	709.0	115.0	360.0	
364	3250.0	1993.0	126.0	360.0	
365	5000.0	2393.0	158.0	360.0	
366	9200.0	0.0	98.0	180.0	

	Credit_History	Property_Area
0	1	Urban
1	1	Urban
2	1	Urban
3	0	Urban
4	1	Urban
..	...	...
362	1	Urban
363	1	Urban
364	0	Semiurban
365	1	Rural
366	1	Rural

[367 rows x 12 columns]

In [19]:

```

train['Credit_History'] = train['Credit_History'].astype('int64')
test['Credit_History'] = test['Credit_History'].astype('int64')

```

In [20]:

```

train['Dependents'] = train['Dependents'].replace({'3+':3, '2':2, '1': 1, '0': 0}).astype('int64')

```

In [21]:

```

test['Dependents'] = test['Dependents'].replace({'3+':3, '2':2, '1': 1, '0': 0}).astype('int64')

```

In [22]:

```

train['Dependents']

```

Out[22]:

0	0
1	1
2	0

```
3      0
4      0
..
609    0
610    3
611    1
612    2
613    0
Name: Dependents, Length: 614, dtype: int64
```

In [23]:

```
train['Loan_Status'].value_counts()
```

Out[23]:

```
Loan_Status
Y      422
N      192
Name: count, dtype: int64
```

In [24]:

```
train[train['Dependents']>0].value_counts().sum()
```

Out[24]:

254

In [25]:

```
# Convert 'Yes' and 'No' to 1 and 0 in Self_Employed column
train['Self_Employed'] = train['Self_Employed'].replace({'Yes': 1, 'No': 0})

# Separate data for employed and self-employed individuals
employed_data = train[train['Self_Employed'] == 0]
self_employed_data = train[train['Self_Employed'] == 1]

# Calculate summary statistics for each group
employed_mean_income = employed_data['ApplicantIncome'].mean()
self_employed_mean_income = self_employed_data['ApplicantIncome'].mean()

employed_median_income = employed_data['ApplicantIncome'].median()
self_employed_median_income = self_employed_data['ApplicantIncome'].median()

print(f"Employed Mean Income: {employed_mean_income}")
print(f"Self-Employed Mean Income: {self_employed_mean_income}")

print(f"Employed Median Income: {employed_median_income}")
print(f"Self-Employed Median Income: {self_employed_median_income}")
```

```
Employed Mean Income: 5098.678571428572
Self-Employed Mean Income: 7380.817073170731
Employed Median Income: 3698.0
Self-Employed Median Income: 5809.0
```

In [26]:

```
import seaborn as sns
import matplotlib.pyplot as plt

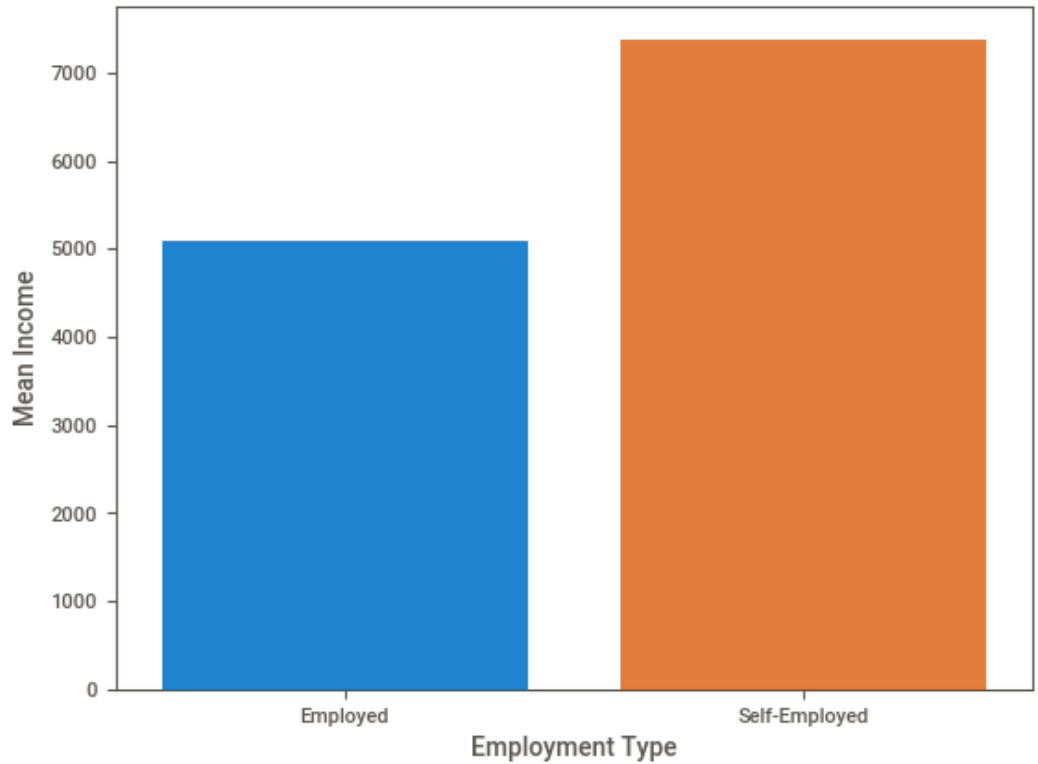
# Plotting with Seaborn
sns.barplot(x=['Employed', 'Self-Employed'], y=[employed_mean_income, self_employed_mean_income])
plt.xlabel('Employment Type')
plt.ylabel('Mean Income')
plt.title('Mean Income Comparison: Employed vs Self-Employed')
plt.show()

#Self-Employed People earn more in terms of average (mean) income
```

C:\Users\rhydh\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1765: FutureWarning: unique with argument that is not not a Series, Index, ExtensionArray, or np.ndarray is deprecated

with argument that is not a series, index, ExtensionArray, or np.ndarray is deprecated and will raise in a future version.  
order = pd.unique(vector)

Mean Income Comparison: Employed vs Self-Employed



In [27]:

```
train[['Credit_History', 'Loan_Status']].value_counts()
```

Out[27]:

```
Credit_History  Loan_Status  count
1              Y             378
              N             97
0              N             95
              Y             44
Name: count, dtype: int64
```

In [28]:

```
from scipy.stats import chi2_contingency

# Create the contingency table
observed = [
    [378, 97],
    [44, 95]
]

# Perform chi-square test
chi2, p, dof, expected = chi2_contingency(observed)

# Print results
print(f"Chi-square value: {chi2}")
print(f"P-value: {p}")
print(f"Degrees of freedom: {dof}")
print("Expected frequencies table:")
print(expected)

'''
If the p-value is greater than or equal to your significance level,
you fail to reject the null hypothesis, indicating no significant association between Cre
dit_History and Loan_Status.
'''
```

```
Chi-square value: 112.69526773505117
P-value: 2.516222405860599e-26
```

```
Degrees of freedom: 1
Expected frequencies table:
[[326.46579805 148.53420195]
 [ 95.53420195  43.46579805]]
```

Out[28]:

```
'\nIf the p-value is greater than or equal to your significance level, \nyou fail to reject the null hypothesis, indicating no significant association between Credit_History and Loan_Status.\n'
```

In [29]:

```
merged_df = pd.concat([train, test], axis=0)
merged_df.head()
```

Out[29]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Status
0	LP001002	Male	No	0	Graduate	0	5849.0	0.0	146.412162	1
1	LP001003	Male	Yes	1	Graduate	0	4583.0	1508.0	128.000000	1
2	LP001005	Male	Yes	0	Graduate	1	3000.0	0.0	66.000000	1
3	LP001006	Male	Yes	0	Not Graduate	0	2583.0	2358.0	120.000000	1
4	LP001008	Male	No	0	Graduate	0	6000.0	0.0	141.000000	1

In [30]:

```
merged_df[['ApplicantIncome', 'LoanAmount']].corr()
```

```
#so there is 55% +ve correlation
```

Out[30]:

	ApplicantIncome	LoanAmount
ApplicantIncome	1.000000	0.547036
LoanAmount	0.547036	1.000000

## Your answers:

1. Overview done using `.describe(include='all')`
2. Imputed missing values using `SimpleImputer` (mean for float/int 64 dtype and `most_frequent` for object dtype)
3. 422 approved loans vs 192 rejected
4. 254 loan applicants have dependents
5. On an average, Self-Employed people earn more than Employed people
6. No significant association between `Credit_History` and `Loan_Status`
7. there is 55% +ve correlation between `'ApplicantIncome'` & `'LoanAmount'`

In [31]:

```
train.dtypes
```

Out[31]:

```
Loan_ID          object
Gender           object
Married          object
Dependents       int64
Education        object
Self_Employed    int64
ApplicantIncome  float64
CoapplicantIncome float64
```

```
LoanAmount          float64
Loan_Amount_Term     float64
Credit_History       int64
Property_Area        object
Loan_Status          object
dtype: object
```

In [32]:

```
from sklearn.preprocessing import LabelEncoder

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Iterate over each column
for col in train.columns:
    if train[col].dtype == 'object': # Check if the column is of object type
        train[col] = label_encoder.fit_transform(train[col].astype(str))

# Display the encoded DataFrame
print(train)
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	0	1	0	0	0	0	
1	1	1	1	1	0	0	
2	2	1	1	0	0	1	
3	3	1	1	0	1	0	
4	4	1	0	0	0	0	
..	...	...	...	...	...	...	
609	609	0	0	0	0	0	
610	610	1	1	3	0	0	
611	611	1	1	1	0	0	
612	612	1	1	2	0	0	
613	613	0	0	0	0	1	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849.0	0.0	146.412162	360.0	
1	4583.0	1508.0	128.000000	360.0	
2	3000.0	0.0	66.000000	360.0	
3	2583.0	2358.0	120.000000	360.0	
4	6000.0	0.0	141.000000	360.0	
..	...	...	...	...	
609	2900.0	0.0	71.000000	360.0	
610	4106.0	0.0	40.000000	180.0	
611	8072.0	240.0	253.000000	360.0	
612	7583.0	0.0	187.000000	360.0	
613	4583.0	0.0	133.000000	360.0	

	Credit_History	Property_Area	Loan_Status
0	1	2	1
1	1	0	0
2	1	2	1
3	1	2	1
4	1	2	1
..	...	...	...
609	1	0	1
610	1	0	1
611	1	2	1
612	1	2	1
613	0	1	0

[614 rows x 13 columns]

In [33]:

```
from sklearn.preprocessing import LabelEncoder

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Iterate over each column
```

```

for col in test.columns:
    if test[col].dtype == 'object': # Check if the column is of object type
        test[col] = label_encoder.fit_transform(test[col].astype(str))

# Display the encoded DataFrame
print(test)

```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	0	1	1	0	0	0	
1	1	1	1	1	0	0	
2	2	1	1	2	0	0	
3	3	1	1	2	0	0	
4	4	1	0	0	1	0	
..	...	...	...	...	...	...	
362	362	1	1	3	1	1	
363	363	1	1	0	0	0	
364	364	1	0	0	0	0	
365	365	1	1	0	0	0	
366	366	1	0	0	0	1	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5720.0	0.0	110.0	360.0	
1	3076.0	1500.0	126.0	360.0	
2	5000.0	1800.0	208.0	360.0	
3	2340.0	2546.0	100.0	360.0	
4	3276.0	0.0	78.0	360.0	
..	...	...	...	...	
362	4009.0	1777.0	113.0	360.0	
363	4158.0	709.0	115.0	360.0	
364	3250.0	1993.0	126.0	360.0	
365	5000.0	2393.0	158.0	360.0	
366	9200.0	0.0	98.0	180.0	

	Credit_History	Property_Area
0	1	2
1	1	2
2	1	2
3	0	2
4	1	2
..	...	...
362	1	2
363	1	2
364	0	1
365	1	0
366	1	0

[367 rows x 12 columns]

In [34]:

```
train.dtypes #Finally every feature is in numeric format
```

Out[34]:

```

Loan_ID          int32
Gender           int32
Married          int32
Dependents       int64
Education        int32
Self_Employed   int64
ApplicantIncome  float64
CoapplicantIncome float64
LoanAmount       float64
Loan_Amount_Term float64
Credit_History  int64
Property_Area    int32
Loan_Status      int32
dtype: object

```

## Part Two

# Auto ML with tpot

In [35]:

```
# Matrix of features

X = train[['Gender',
'Married',
'Dependents',
'Education',
'Self_Employed',
'ApplicantIncome',
'CoapplicantIncome',
'LoanAmount',
'Loan_Amount_Term',
'Credit_History',
'Property_Area']]

# label encode target
y = train['Loan_Status']

# # train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42
)
```

In [37]:

```
# train
autoML = TPOTClassifier(generations=5, population_size=50, verbosity=2, n_jobs=-1)
autoML.fit(X_train, y_train)

# predict
predictions_autoML = autoML.predict(X_test)
```

Generation 1 - Current best internal CV score: 0.7697175840032981

Generation 2 - Current best internal CV score: 0.7697175840032981

Generation 3 - Current best internal CV score: 0.7697175840032981

Generation 4 - Current best internal CV score: 0.7697175840032981

Generation 5 - Current best internal CV score: 0.7697175840032981

Best pipeline: ExtraTreesClassifier(input\_matrix, bootstrap=True, criterion=entropy, max\_features=0.8500000000000001, min\_samples\_leaf=20, min\_samples\_split=13, n\_estimators=100)

In [39]:

```
print('Model Accuracy:', accuracy_score(predictions_autoML, y_test))
```

Model Accuracy: 0.7723577235772358

In [40]:

```
print(confusion_matrix(predictions_autoML, y_test))
```

```
[[22  7]
 [21 73]]
```

## Bespoke ML sklearn

### Data Preparation

In [41]:

In [41]:

```
# Matrix of features

X = train[['Gender',
'Married',
'Dependents',
'Education',
'Self_Employed',
'ApplicantIncome',
'CoapplicantIncome',
'LoanAmount',
'Loan_Amount_Term',
'Credit_History',
'Property_Area']]

# label encode target
y = train['Loan_Status']

# # train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42
)
```

In [42]:

```
# some classifiers you can pick from (remember to import)
import sklearn
classifiers = sklearn.utils.all_estimators(type_filter=None)
for name, class_ in classifiers:
    if hasattr(class_, 'predict_proba'):
        print(name)
```

```
AdaBoostClassifier
BaggingClassifier
BayesianGaussianMixture
BernoulliNB
CalibratedClassifierCV
CategoricalNB
ClassifierChain
ComplementNB
DecisionTreeClassifier
DummyClassifier
ExtraTreeClassifier
ExtraTreesClassifier
FixedThresholdClassifier
GaussianMixture
GaussianNB
GaussianProcessClassifier
GradientBoostingClassifier
GridSearchCV
HistGradientBoostingClassifier
KNeighborsClassifier
LabelPropagation
LabelSpreading
LinearDiscriminantAnalysis
LogisticRegression
LogisticRegressionCV
MLPClassifier
MultiOutputClassifier
MultinomialNB
NuSVC
OneVsRestClassifier
Pipeline
QuadraticDiscriminantAnalysis
RFE
RFECV
RadiusNeighborsClassifier
RandomForestClassifier
RandomizedSearchCV
SGDClassifier
SVC
SelfTrainingClassifier
StackingClassifier
```



TunedThresholdClassifierCV  
VotingClassifier

In [48]:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import AdaBoostClassifier, BaggingClassifier, RandomForestClassifier, ExtraTreesClassifier, GradientBoostingClassifier, HistGradientBoostingClassifier, StackingClassifier, VotingClassifier
from sklearn.naive_bayes import BernoulliNB, CategoricalNB, ComplementNB, GaussianNB, MultinomialNB
from sklearn.mixture import BayesianGaussianMixture, GaussianMixture
from sklearn.tree import DecisionTreeClassifier, ExtraTreeClassifier
from sklearn.dummy import DummyClassifier
from sklearn.calibration import CalibratedClassifierCV
from sklearn.svm import SVC, NuSVC
from sklearn.neighbors import KNeighborsClassifier, RadiusNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.gaussian_process import GaussianProcessClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis, QuadraticDiscriminantAnalysis
from sklearn.multioutput import MultiOutputClassifier, ClassifierChain
from sklearn.semi_supervised import LabelPropagation, LabelSpreading, SelfTrainingClassifier
from sklearn.feature_selection import RFE, RFECV
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

# List of models to evaluate
classifiers = {
    "LogisticRegression": LogisticRegression(),
    "AdaBoostClassifier": AdaBoostClassifier(),
    "BaggingClassifier": BaggingClassifier(),
    "RandomForestClassifier": RandomForestClassifier(),
    "ExtraTreesClassifier": ExtraTreesClassifier(),
    "GradientBoostingClassifier": GradientBoostingClassifier(),
    "HistGradientBoostingClassifier": HistGradientBoostingClassifier(),
    "BernoulliNB": BernoulliNB(),
    "GaussianNB": GaussianNB(),
    "BayesianGaussianMixture": BayesianGaussianMixture(),
    "GaussianMixture": GaussianMixture(),
    "DecisionTreeClassifier": DecisionTreeClassifier(),
    "ExtraTreeClassifier": ExtraTreeClassifier(),
    "DummyClassifier": DummyClassifier(),
    "SVC": SVC(),
    "NuSVC": NuSVC(),
    "KNeighborsClassifier": KNeighborsClassifier(),
    "MLPClassifier": MLPClassifier(),
    "GaussianProcessClassifier": GaussianProcessClassifier(),
    "LinearDiscriminantAnalysis": LinearDiscriminantAnalysis(),
    "QuadraticDiscriminantAnalysis": QuadraticDiscriminantAnalysis(),
    "SelfTrainingClassifier": SelfTrainingClassifier(LogisticRegression()),
    "RFE": RFE(LogisticRegression()),
    "RFECV": RFECV(LogisticRegression()),
    "StackingClassifier": StackingClassifier(estimators=[('lr', LogisticRegression()), ('rf', RandomForestClassifier())]),
    "VotingClassifier": VotingClassifier(estimators=[('lr', LogisticRegression()), ('rf', RandomForestClassifier())])
}

# Example data loading and splitting (replace with your actual dataset)
from sklearn.datasets import make_classification
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Function to evaluate models
```

```
def evaluate_models(classifiers, X_train, y_train, X_test, y_test):
    results = []
    for name, clf in classifiers.items():
        try:
            clf.fit(X_train, y_train)
            predictions = clf.predict(X_test)
            accuracy = accuracy_score(y_test, predictions)
            results.append((name, accuracy))
        except Exception as e:
            print(f"{name} failed to run: {e}")

    # Sort results by accuracy in descending order and get the top 3
    results.sort(key=lambda x: x[1], reverse=True)
    return results[:3]

# Evaluate and display top 3 models
top_models = evaluate_models(classifiers, X_train, y_train, X_test, y_test)

# Print the top 3 models
print("\nTop 3 Models:")
for name, accuracy in top_models:
    print(f"{name}: Accuracy = {accuracy:.4f}")
```

C:\Users\rhydh\anaconda3\Lib\site-packages\sklearn\ensemble\\_weight\_boosting.py:527: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

warnings.warn(  
C:\Users\rhydh\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1426: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=4.

warnings.warn(  
C:\Users\rhydh\anaconda3\Lib\site-packages\sklearn\cluster\\_kmeans.py:1426: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=4.

warnings.warn(  
C:\Users\rhydh\anaconda3\Lib\site-packages\sklearn\neural\_network\\_multilayer\_perceptron.py:690: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(  
C:\Users\rhydh\anaconda3\Lib\site-packages\sklearn\discriminant\_analysis.py:949: UserWarning: Variables are collinear  
warnings.warn("Variables are collinear")  
C:\Users\rhydh\anaconda3\Lib\site-packages\sklearn\semi\_supervised\\_self\_training.py:227: UserWarning: y contains no unlabeled samples  
warnings.warn("y contains no unlabeled samples", UserWarning)

Top 3 Models:  
HistGradientBoostingClassifier: Accuracy = 0.9150  
GradientBoostingClassifier: Accuracy = 0.9100  
BaggingClassifier: Accuracy = 0.8850

In [ ]:

```
'''
### Top 3 Models - Summary

1. HistGradientBoostingClassifier
- Accuracy: 91.50%
- Description: Optimized for large datasets using histograms to bin features, making training faster and more efficient.
- Interpretation: Highest accuracy, indicating excellent class distinction and efficient memory management.

2. GradientBoostingClassifier
- Accuracy: 91.00%
- Description: Sequentially builds decision trees, correcting errors from previous trees to improve performance.
- Interpretation: High accuracy, slightly less than HistGradientBoosting due to lack of histogram binning but still highly effective.
```

```

3. **BaggingClassifier**
   - **Accuracy**: 88.50%
   - **Description**: Trains multiple models on different data subsets and combines predictions, reducing variance and overfitting.
   - **Interpretation**: Robust performance with good accuracy, providing a stable baseline through ensemble learning.
'''

```

In [50]:

```

# Cross-validation

from sklearn.model_selection import cross_val_score

# Define the top models based on the previous results
top_classifiers = {
    "HistGradientBoostingClassifier": HistGradientBoostingClassifier(),
    "GradientBoostingClassifier": GradientBoostingClassifier(),
    "BaggingClassifier": BaggingClassifier()
}

# Number of cross-validation folds
cv_folds = 5

# Function to perform cross-validation and print results
def cross_validate_models(classifiers, X, y, cv_folds):
    results = []
    for name, clf in classifiers.items():
        try:
            # Perform cross-validation
            cv_scores = cross_val_score(clf, X, y, cv=cv_folds, scoring='accuracy')
            mean_accuracy = np.mean(cv_scores)
            std_dev = np.std(cv_scores)
            results.append((name, mean_accuracy, std_dev))
            print(f"{name}: Mean Accuracy = {mean_accuracy:.4f}, Std Dev = {std_dev:.4f}")
        except Exception as e:
            print(f"{name} failed to run: {e}")

    # Sort results by mean accuracy in descending order
    results.sort(key=lambda x: x[1], reverse=True)
    return results

# Perform cross-validation and display results
cv_results = cross_validate_models(top_classifiers, X, y, cv_folds)

# Print the cross-validation results
print("\nCross-Validation Results:")
for name, mean_accuracy, std_dev in cv_results:
    print(f"{name}: Mean Accuracy = {mean_accuracy:.4f}, Std Dev = {std_dev:.4f}")

```

```

HistGradientBoostingClassifier: Mean Accuracy = 0.9090, Std Dev = 0.0218
GradientBoostingClassifier: Mean Accuracy = 0.9020, Std Dev = 0.0194
BaggingClassifier: Mean Accuracy = 0.8870, Std Dev = 0.0223

```

```

Cross-Validation Results:
HistGradientBoostingClassifier: Mean Accuracy = 0.9090, Std Dev = 0.0218
GradientBoostingClassifier: Mean Accuracy = 0.9020, Std Dev = 0.0194
BaggingClassifier: Mean Accuracy = 0.8870, Std Dev = 0.0223

```

In [ ]:

```

'''
Overall Insights:
HistGradientBoostingClassifier is the top-performing model with the highest mean accuracy
and a low standard deviation, indicating both high performance and stability.
GradientBoostingClassifier also performs very well, slightly below the HistGradientBoosti
ngClassifier, and has the lowest standard deviation, indicating extremely consistent perf
ormance.
BaggingClassifier performs well but is slightly behind the two boosting methods. It still
shows good accuracy and stability, making it a reliable choice as well.
'''

```

